

OpenMP Tasks



Oregon State
University
Mike Bailey


mjb@cs.oregonstate.edu



This work is licensed under a [Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International License](https://creativecommons.org/licenses/by-nc-nd/4.0/)



Oregon State
University
Computer Graphics



tasks.pptx

mjb - March 23, 2020

Remember OpenMP Sections?

Sections are independent blocks of code, able to be assigned to separate threads if they are available.

```
#pragma omp parallel sections
{
    #pragma omp section
    {
        Task 1
    }
    #pragma omp section
    {
        Task 2
    }
}
```

There is an **implied barrier** at the end



Oregon State
University
Computer Graphics

OpenMP sections are **static**, that is, they are good if you know, when you are writing the program, how many of them you will need.

mjb - March 23, 2020

It would be nice to have something more Dynamic

3



Imagine a capability where you can write something to do down on a Post-It® note, accumulate the Post-It notes, then have all of the threads together execute that set of tasks.

You would also like to not have to know, ahead of time, how many of these Post-It notes you will write. That is, you want the total number to be **dynamic**.

Well, congratulations, you have just invented **OpenMP Tasks**!

Oregon State University
Computer Graphics

mjb - March 23, 2020

OpenMP Tasks

4

- An OpenMP task is a single line of code or a structured block which is immediately “written down” in a list of tasks.
- The new task can be executed immediately, or it can be deferred.
- If the *if* clause is used and the argument evaluates to 0, then the task is executed immediately, superseding whatever else that thread is doing.
- There has to be an existing parallel thread team for this to work. Otherwise one thread ends up doing all tasks and you don’t get any contribution to parallelism.
- One of the best uses of this is to process elements of a linked list or a tree.

You can create a task barrier with:

#pragma omp taskwait

Tasks are very much like OpenMP **Sections**, but Sections are static, that is, the number of sections is set when you write the code, whereas **Tasks** can be created anytime, and in any number, under control of your program’s logic.


Oregon State University
Computer Graphics

mjb - March 23, 2020

OpenMP Task Example: Something (Supposedly) Simple

5


```
omp_set_num_threads( 2 );
#pragma omp parallel default(none)
{
    #pragma omp task
    fprintf( stderr, "A\n" );

    #pragma omp task
    fprintf( stderr, "B\n" );
}
```

Without this, thread #0 has to do everything

Writes fprintf(stderr, "A\n"); on a sticky note and adds it to the list

Writes fprintf(stderr, "B\n"); on a sticky note and adds it to the list



Oregon State
University
Computer Graphics

mjb - March 23, 2020

If You Run This a Number of Times, You Get This: (What Happened?)

6

Run #	1	2	3	4	5
	B	B	B	B	B
	A	B	A	A	A
	B	A	A	A	B
	A	A	B	B	A

1. Why do we not get the same output every time?
2. Why do we get 4 things printed when we only have print statements in 2 tasks?

Not so simple, huh?

The first answer is easy. Unless you make some special arrangements, the order of execution of the different tasks is *undefined*.

The second answer is that we actually asked each of the two threads to put two tasks on the sticky notes, for a total of four. How can we get only one thread to do this?



Oregon State
University
Computer Graphics


mjb - March 23, 2020

The “single” Pragma

7

```
omp_set_num_threads( 2 );
#pragma omp parallel default(none)
{
    #pragma omp single
    {
        #pragma omp task
        fprintf( stderr, "A\n" );

        #pragma omp task
        fprintf( stderr, "B\n" );
    }
}
```



When using Tasks, you only want *one* thread to write the things to do down on the sticky note, but you want *all* of the threads to be able to execute the sticky notes.



Oregon State
University
Computer Graphics

mjb - March 23, 2020

But, if you run this, the order of printing will still be non-deterministic. To solve that problem, do this:


8

```
omp_set_num_threads( 2 );
#pragma omp parallel
{
    #pragma omp single default(none)
    {
        #pragma omp task
        fprintf( stderr, "A\n" );

        #pragma omp taskwait ← Causes all tasks to wait until they are completed

        #pragma omp task
        fprintf( stderr, "B\n" );

        #pragma omp taskwait ← Causes all tasks to wait until they are completed
    }
}
```



Oregon State
University
Computer Graphics

mjb - March 23, 2020

A Better OpenMP Task Example: Processing each Element of a Linked List

9

Without this, thread #0 has to do everything

```
#pragma omp parallel default(none)
{
    Without this, each thread does a full traversal – bad idea!

    #pragma omp single default(none)
    {
        element *p = listHead;
        while( p != NULL )
        {
            Write "Process( p )" on a sticky note and add it to the list

            #pragma omp task firstprivate(p)
            Process( p );

            p = p->next;
        }

        #pragma omp taskwait
    }
}
```

Writes the current value of p into the task and immediately makes it private (i.e., not shared)

Put this here if you want to wait for all tasks to finish being executed before proceeding

University
Computer Graphics

mjb – March 23, 2020

One more thing – Task Dependencies

10

Remember from before: unless you make some special arrangements, the order of execution of the different tasks is *undefined*. Here come the special arrangements.

```
omp_set_num_threads( 3 );
#pragma omp parallel
{
    #pragma omp single default(none)
    {
        float a, b, c;
        #pragma omp task depend( OUT: a )
        a = 10.;

        #pragma omp task depend( IN: a, OUT: b )
        b = a + 16.;

        #pragma omp task depend( IN: b )
        c = b + 12.;
    }

    #pragma omp taskwait
}
```



This maintains the proper dependencies, but, because it involves all of the tasks, it essentially serializes the parallelism out of them.
Be careful not to go overboard with dependencies!




mjb – March 23, 2020

Tree Traversal Algorithms

11

Given a tree:



- We would like to traverse it as quickly as possible.
- We are assuming that we do not need to traverse it in order.
- We just need to visit all nodes.

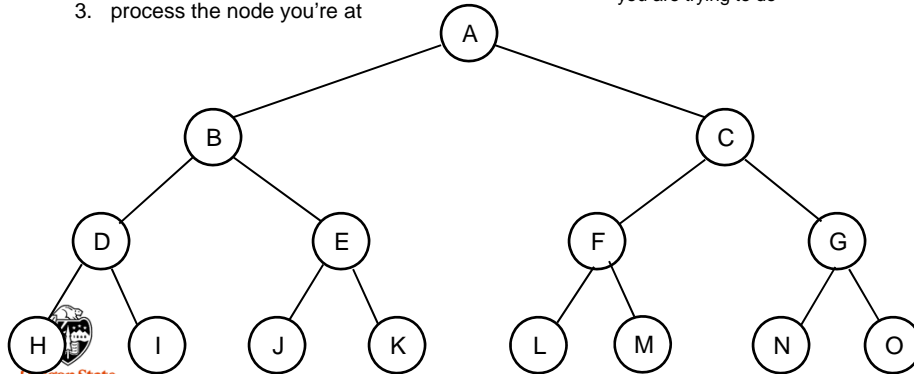
mjb - March 23, 2020

Tree Traversal Algorithms

12

- This is common in graph algorithms, such as searching.
- If the tree is binary and is balanced, then the maximum depth of the tree is $\log_2(\text{\# of Nodes})$
- Strategy at a node:
 1. follow one descendent node
 2. follow the other descendent node
 3. process the node you're at

This order could be re-arranged, depending on what you are trying to do



mjb - March 23, 2020

Tree Traversal Algorithms

13



```
#pragma omp parallel
```

```
#pragma omp single
```

```
Traverse( root );
```

```
#pragma omp taskwait
```

Without this, thread #0 has to do everything

Without this, each thread does a full traversal – bad idea!

Put this here if you want to wait for all nodes to be traversed before proceeding

Parallelizing a Binary Tree Traversal with Tasks

14



```
void
Traverse( Node *n )
{
    if( n->left != NULL )
    {
        #pragma omp task private(n) untied
        Traverse( n->left );
    }

    if( n->right != NULL )
    {
        #pragma omp task private(n) untied
        Traverse( n->right );
    }

    #pragma omp taskwait


    Process( n );
}
```

Put this here if you want to wait for both branches to be taken before processing the parent

Parallelizing a Binary Tree Traversal with Tasks

15

Traverse(A);

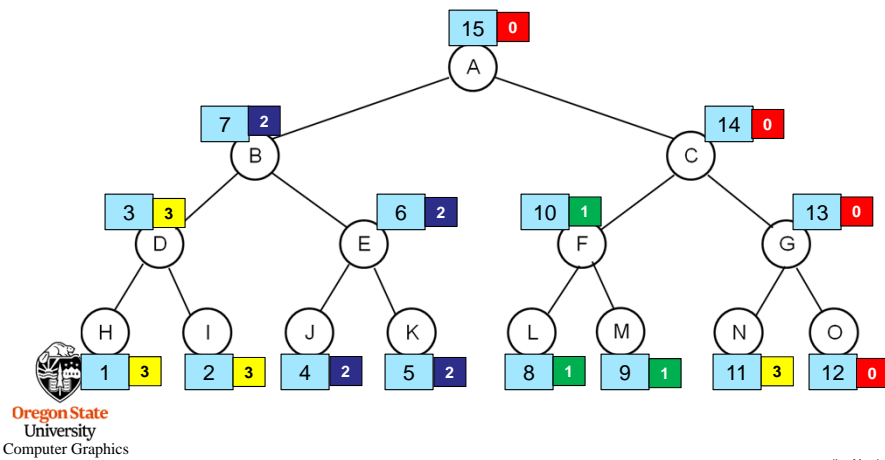


Parallelizing a Binary Tree Traversal with Tasks: *Tied* (gcc 8.2)

16

Threads:
0 1 2 3

Traverse(A);

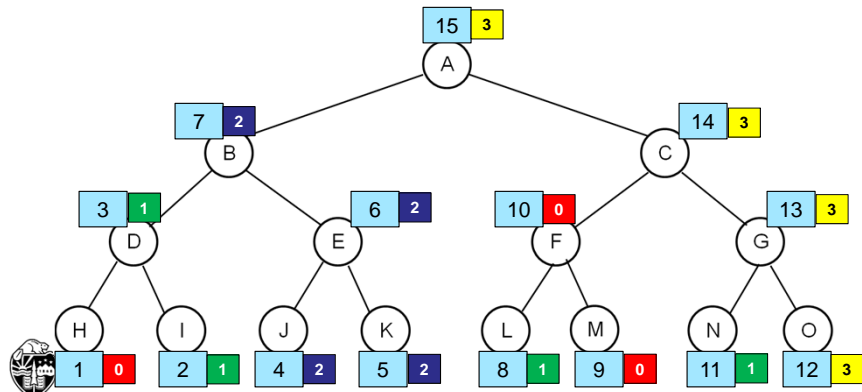


Parallelizing a Binary Tree Traversal with Tasks: *Untied* (gcc 8.2)

17

Threads:
0 1 2 3

Traverse(A);



Oregon State
University
Computer Graphics

mjb - March 23, 2020

How Evenly Tasks Get Assigned to Threads

18

6 Levels – g++ 4.9:

Thread #	Number of Tasks
0	1
1	32
2	47
3	47

6 Levels – g++ 8.2:

Thread #	Number of Tasks
0	31
1	46
2	32
3	18

12 Levels – g++ 4.9:

Thread #	Number of Tasks
0	2561
1	2
2	2813
3	2815

12 Levels – g++ 8.2:

Thread #	Number of Tasks
0	1
1	2048
2	3071
3	3071

Oregon State
University
Computer Graphics

mjb - March 23, 2020

How Evenly Tasks Get Assigned to Threads

19

6 Levels – g++ 8.2:

Thread #	Number of Tasks
0	31
1	46
2	32
3	18

6 Levels – icpc 15.0.0:

Thread #	Number of Tasks
0	29
1	31
2	41
3	26

12 Levels – g++ 8.2:

Thread #	Number of Tasks
0	1
1	2048
2	3071
3	3071

12 Levels – icpc 15.0.0:

Thread #	Number of Tasks
0	1999
1	2068
2	2035
3	2089



Oregon State
University
Computer Graphics

mjb – March 23, 2020

Benchmarking a Binary Task-driven Tree Traversal

20

```
void
Process( Node *n )
{
    for( int i = 0; i < 1024; i++ )
    {
        n->value = pow( n->value, 1.1 );
    }
}
```




Oregon State
University
Computer Graphics

mjb – March 23, 2020


Performance vs. Number of Threads

21



Performance vs. Number of Levels

22



Parallelizing a Tree Traversal with Tasks

23

- Tasks get spread among the current “thread team”
- Tasks can execute immediately or can be deferred. They are executed at “some time”.
- Tasks can be moved between threads, that is, if one thread has a backlog of tasks to do, an idle thread can come steal some workload.
- Tasks are more dynamic than sections. The task paradigm would still work if there was a variable number of children at each node.

Parallelizing an N-Tree Traversal with Tasks

24

```
void
Traverse( Node *n )
{
    for( int i = 0; i < n->numChildren; i++ )
    {
        if( n->child[ i ] != NULL )
        {
            #pragma omp task
            Traverse( n->child[ i ] );
        }
    }

    #pragma omp taskwait

    Process( n );
}
```

